

The Art of Fuzzing Without Fuzzing



@rantyben

Kiwicon 2015

Fuzzing (A History)

The year is 2002...

Buffy
the vampire slayer

spike

visit the official buffy site at <http://www.buffy.com>





Whitebox, Greybox, Polka Dots

- **SAGE** Godefroid / MS Research
- **EFS** DeMott
- **Sidewinder** Embleton, Sparks, Cunningham
- **Bunny the Fuzzer** Zalewski
- **Flayer** Ormandy
- **KLEE** Cadar et al
- ... many more

Whitebox, Greybox, Polka Dots

1. Instrument Stuff
2. Do things
3. Measure
4. Do better things

Problems

- State Explosion
- Devil In Details
- Performance
- Limited Availability
- Fragile Code
- Artificial Results
- Require Source





Charlie Miller



The “Army of Monkeys” Talk

- CSW 2010
- Collected ~80k PDFs
- Traced for coverage, reduced to ~1500
- Fuzzed several parsers with “millerfuzz”
- ~7 million iterations TOTAL (!!)
- Drowned in bugs



Prospector / BM3

- Me, 2010
- Scraper using Bing
- Tracer using Pin
- Distributed VM based fuzz cluster
- ~200 VMs, peak speeds ~50t/s
- Scales to 100s of Mt, works on Windows

Problems

- Scraping slow
- Office is slow
- VMs are slow
- “Real” files don’t exercise much code
- Millerfuzz is a pretty bad mutator
- Bug probability is constant



“Block” Based Fuzzing?

- Invented by Archimedes ~250BC
- Acts like Gears + Ratchet
- One Human easily lifts 1-2 tons
- Force Multiplier + Hill Climbing

roar.



AFL - It's What's Up.

- Instrument source or binary
- Dumb fuzz, get coverage
- Tests that add coverage are kept
- Naturally evolves towards max coverage
- Can appear terrifyingly intelligent

AFL - It's What's Up.

- Novel coverage approach (bitmap)
- VERY VERY VERY FAST
- Well tested dumb mutators
- Zero configuration
- No academic silliness
- Breaks \$#\$%*@* Everything

american fuzzy lop 1.92b (tif-S0)

process timing		overall results	
run time : 9 days, 20 hrs, 29 min, 0 sec		cycles done : 31	
last new path : 0 days, 0 hrs, 2 min, 33 sec		total paths : 9550	
last uniq crash : 1 days, 1 hrs, 5 min, 20 sec		uniq crashes : 205	
last uniq hang : 9 days, 13 hrs, 58 min, 39 sec		uniq hangs : 500+	
cycle progress		map coverage	
now processing : 2885* (30.21%)		map density : 8590 (13.11%)	
paths timed out : 0 (0.00%)		count coverage : 5.37 bits/tuple	
stage progress		findings in depth	
now trying : havoc		favored paths : 776 (8.13%)	
stage execs : 13.2k/20.0k (66.12%)		new edges on : 1349 (14.13%)	
total execs : 711M		total crashes : 4383 (205 unique)	
exec speed : 1269/sec		total hangs : 806k (500+ unique)	
fuzzing strategy yields		path geometry	
bit flips : n/a, n/a, n/a		levels : 33	
byte flips : n/a, n/a, n/a		pending : 3526	
arithmetics : n/a, n/a, n/a		pend fav : 0	
known ints : n/a, n/a, n/a		own finds : 9549	
dictionary : n/a, n/a, n/a		imported : 0	
havoc : 4606/275M, 5148/433M		variable : 1	
trim : 11.36%/1.51M, n/a			

[cpu:126%]

Welcome to the Gt Era

- < 10Mt is now insignificant
- 1Gt MINIMUM per parser
- < 24 hrs with a decent box
- ...but Reader is still slow!



Corpus Driven Fuzzing

- MS15-024 / MS15-029 (Icamtuf)
- IE bugs handling JXR and PNG
- Found without ever fuzzing IE!

Corpus Driven Fuzzing

- Really good corpora are **AMAZING**
- AFL creates really good corpora
- Quickly*
- With Zero* Configuration
- Bonus - they're re-useable!

Building Corpora with AFL

- NOT aimed at finding bugs
- Crashy targets are actually WORSE
- Focus is on exercising code / spec



AFL “Crash” Course

1. Build AFL
2. Build target with AFL compilers
3. Run AFL on instrumented target

Target Selection Tips

- Multiple engines
- Simple apps
 - converters
 - extractors
- Code examples
- Good / Complete code!

The Perfect AFL Target

- Complete protocol support
- Fine-grained error handling
- Exits when done
- Reads STDIN, outputs to STDOUT
- Accepts arbitrary filenames

Instrumentation Tricks

- Patch out
 - extension checking
 - checksum checks
 - disk writes
 - useless config file reads (strace)

Instrumentation Tricks

- Persistent Mode
- Deferred Forkserver (?)
- Speed gains can be HUGE (like 10x)

Persistent Mode - mupdf

```
int main(int argc, char **argv)
{
    char *filename = argc >= 2 ? argv[1] : "";
    render(filename);
    return 0;
}
~
```

108, 1-8

Bot

Persistent Mode - mupdf

```
int main(int argc, char **argv)
{
    char *filename = argc >= 2 ? argv[1] : "";
    while(__AFL_LOOP(100000)) {
        render(filename);
    }
    return 0;
}
```

108,1-8

Bot

Persistent Mode Tips

- Reinitialize ALL dirty variables
- `__AFL_LOOP(xxx)` can limit leaks
- Might even stabilize variable flow!
 - (pdfium example)

AFL Seeds

- Complex formats NEED seeds!
- 1 day x 48 cores with:
 - AFL Sample PDF - 13%
 - Corkami POC PDFS - 30%
 - Add Adobe Samples - 47%

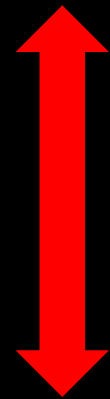
Pollenation

- Targets in one root sync automatically

/path/to/target1/target-S0

/path/to/target1/target-S1

/path/to/target1/target-S2

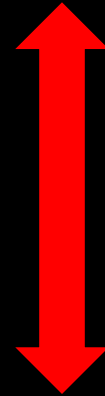


Pollenation

/path/to/target1/t1-S0

/path/to/target1/t1-S1

/path/to/target1/t1-S2



/path/to/target2/t2-S0

/path/to/target2/t2-S1

/path/to/target2/t2-S2



Pollenation

/path/to/target1/t1-S0

/path/to/target1/t1-S1

/path/to/target1/t1-S2

/path/to/target1/t2-S0.sync



/path/to/target2/t2-S0

/path/to/target2/t2-S1

/path/to/target2/t2-S2

/path/to/target2/t1-S0.sync



Pollination

- For core-heavy farms, this is great
- For box-heavy farms try **roving**
- github.com/richoh/roving

“Low Hanging Fruit Trick”

- Run 1 fuzzer with 1 tiny sample
- Run another with varied seeds
- Pollenate after a day
- ~20% smaller total queue
- ~5% faster execution

Spring Cleaning

- AFL is stubborn about resumes
 - Thundering Herd -> timeouts
 - Some queue files 'crash'
- Queue from another target is WORSE

Spring Cleaning

- `export AFL_SKIP_CRASHES=1`
- in `config.h`, try:

```
#define CAL_TMOUT_PERC      200  
#define CAL_TMOUT_ADD      200
```

Spring Cleaning

- Use afl-cmin!
 - Automatically skips crashes
 - Skips timeouts
 - Bonus: actually reduces queue!
- Note: queue/ usually “reinflates” under fuzz

My Software



Doc'or? Why
Carr I thee?



afl-launch

- Launch many afl instances
- Passes through all afl flags
- Redirects stdout to logfiles
- Template filename options

cwtrriage

- Automatically repro and triage crashes
- OSX* and Linux
- Simple, plays well with others
- Output to TXT, JSON or protobuf
- Heavily tested (on the linux side)

cwtriage

- **-tidy** stashes files that don't repro
- **-afl** automatically re-use AFL settings
- **-every** runs on a schedule
- **-file** uses template filenames
 - /dev/shm/xvghktsty.pdf
- **-workers** uses multiple cores

cwdump

- Dumps a cwtrriage database
- Summarises crashes by stack hash



Initiating Demonstration...

```
~ cat > ~/bin/showbugs
```

```
#!/bin/sh
```

```
cwtrriage -root . \  
-workers $(nproc) -af1 > /dev/null
```

```
cwdump ./crashwalk.db > triage.txt
```

```
less triage.txt
```

```
^D
```

```
~ cat > ~/bin/showbugs
```

```
^D
```

```
~ chmod a+x ~/bin/showbugs
```

```
~ cd /my/target && showbugs
```

afl-trivia

- afl-pause / afl-resume
- afl-pcmin
- afl-consolidate
- afl-pollenate

~ cd /rootdir/for/target

~ afl-pause .

~ showbugs

~ afl-resume .

afl-pcmin

- Like afl-cmin, but using GNU parallel
- Minimise a set of files for coverage

afl-pollenate

- Pollenates .sync dirs (shown earlier)
- Runs once an hour, by default

~ afl-pollenate /fuzz/pdf

afl-consolidate

- Consolidate all queue and crash files
- Named via SHA-1 (removes dups)
- Optionally add extension
- Do this to prep for repro stage

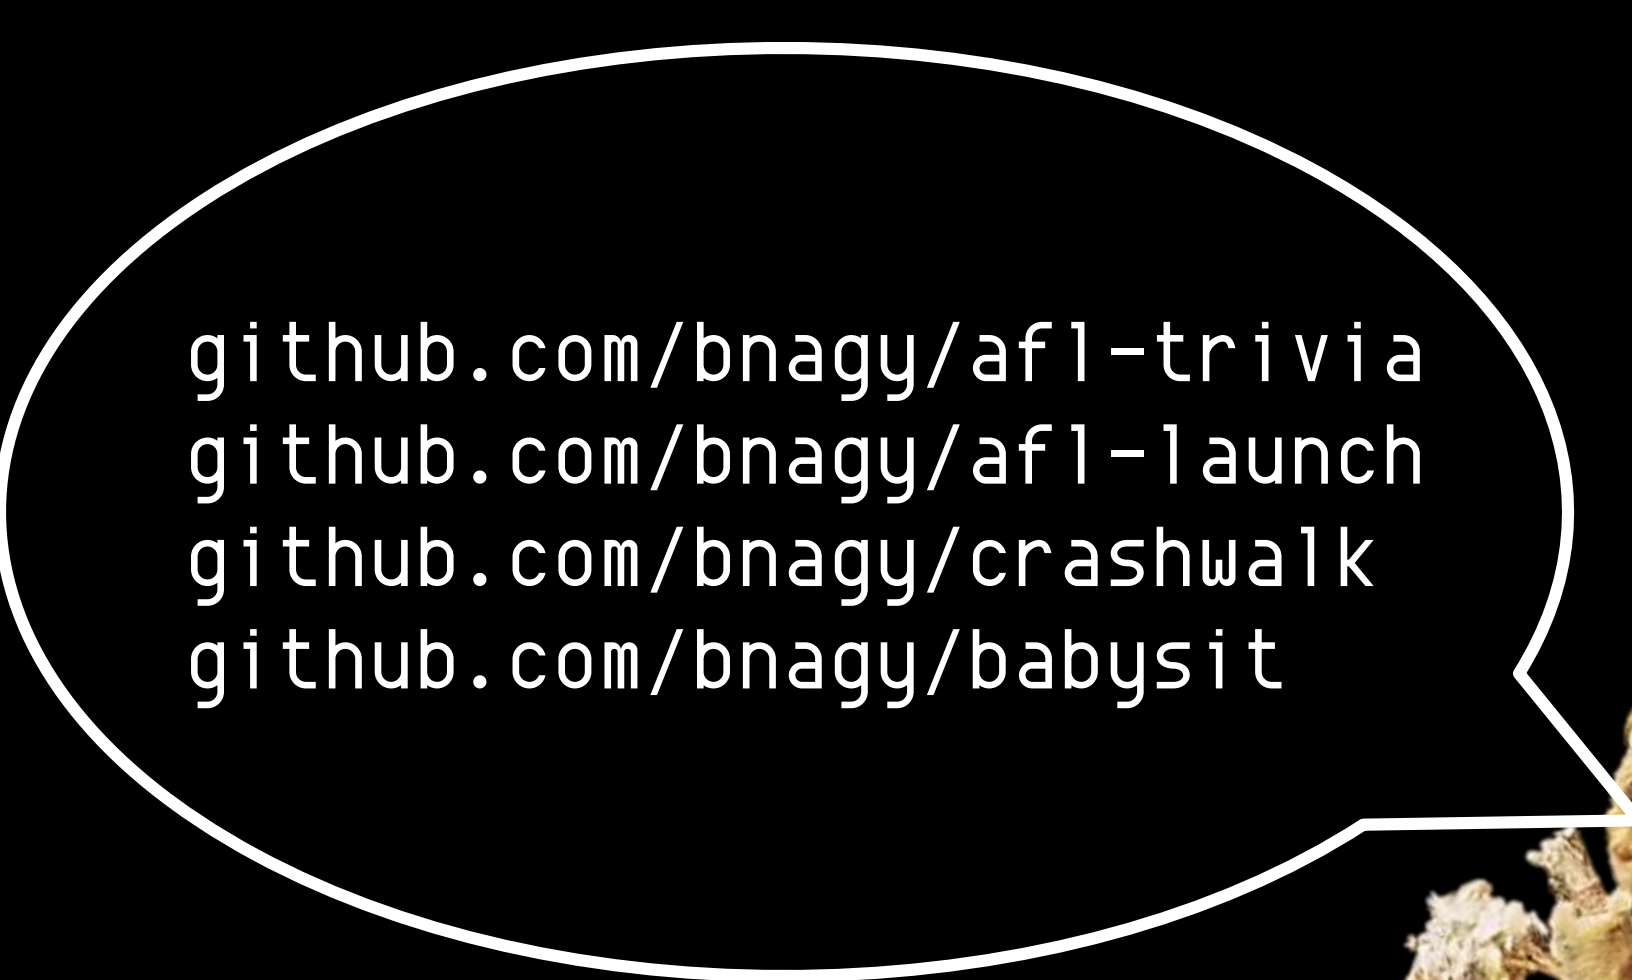
```
~ afl-consolidate . /pdfs .pdf
```

babysit.exe


- Run a Windows target with many files
 - `CreateProcess()`
 - `WaitForSingleObject()`
 - `GetExitCodeProcess()`
- YES it works for /GS violations etc
- BugId is better, but this is like 50x faster
 - github.com/SkyLined/BugId

h0w 2 cyb3rbugs

1. `afl-launch` several linux engines
2. `afl-pollenate` between them
3. `afl-consolidate` that corpus
4. `babysit` real target with those files
5. Use `bugid.py` on tasty exit codes
 - `0xc0000005`
 - `0xc0000409`



github.com/bnagy/afl-trivia
github.com/bnagy/afl-launch
github.com/bnagy/crashwalk
github.com/bnagy/babysit



Ma fathe... I
Carr feer ma
fathe...



Thanks

- Icamtuf for AFL
- Jonathan Foote for GDB exploitable
- ~~@rich0H for the SIGSTOP trick~~
- @snare



I'm feline terrible,
bro! I'm gutted!

**A fuzzer is like
a finger
pointing a way
to the moon.
Don't focus on
the fuzzer, or
you will miss all
that heavenly
glory.**



come at me, bro!

@rantyben

freenode/#afl-users (bnagy)

github.com/bnagy